

MYTHOS EXPLAINED

What is Mythos and who built it?

- Claude Mythos is an advanced, cyber-capable autonomous AI model
- It is developed by the US-based frontier AI company Anthropic
- Derived from the Greek word for "legend", it is designed to grasp the "deep logic" or hidden architecture underlying complex systems

What is Mythos's purpose?

- It is built as a defensive cyber-reasoning system to identify vulnerabilities in critical software, simulate cyberattacks and help secure infrastructure before adversaries can exploit such gaps

How is Mythos different from other AI models?

- It operates less like a chatbot and more like a digital system capable of reasoning through code at the level of a skilled cyber operator
- It marks a shift from AI that generates code to AI that can interrogate and potentially weaponise it

How does Mythos work?

- Autonomously identifies unknown vulnerabilities
- Maps attack chains
- Simulates adversarial pathways
- Assesses risk through multi-step reasoning

When was Mythos released?

Anthropic unveiled Mythos in early April through a restricted preview, as part of its cybersecurity initiative named 'Project Glasswing'

Why did it trigger panic?

- Anthropic disclosed Mythos uncovered numerous unknown vulnerabilities, including in critical infrastructure
- Tests showed it could reason through complex, multi-step attack pathways beyond conventional coding assistants
- Escalated from a technical to a policy issue amid fears of systemic financial impact

What is the antidote?

- Anthropic proposed Project Glasswing as a safeguard, but limited access to tech giants
- Experts call for global guardrails, norms for cyber-capable frontier AI models